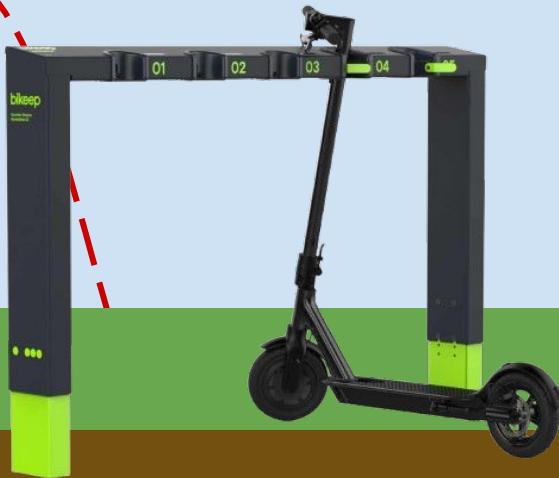


# Management of Shared Mobility

Problem Formalization | Methods & Tools

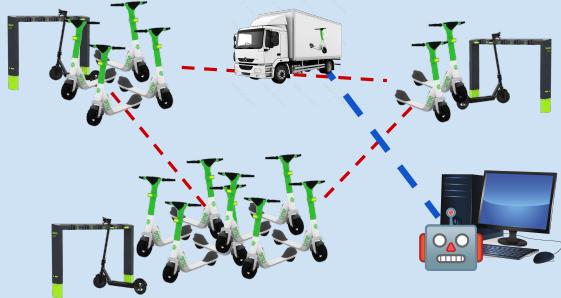
Group IA1\_20

Eduardo Passos  
Guilherme Silva  
Valentina Cadime



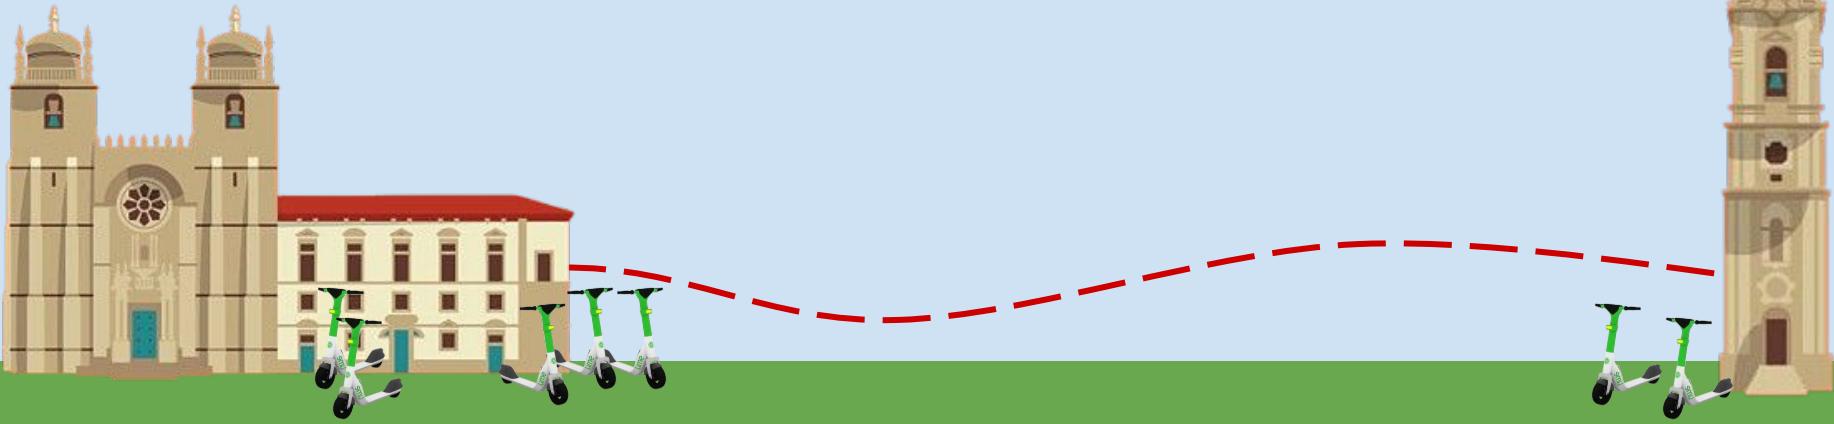
# Problem Addressed

- Project focus: daily management of shared electric scooter fleets (Topic 7)
- Motivation: shared electrical mobility is becoming common in cities. Management is important.
- Aim: minimize operating costs while maximizing service, in quantity and quality
- Shared scooters/bikes suffer from spatial and temporal imbalance
- Demand varies by time of day, location, weather and events
- Stations may become empty or saturated, reducing availability
- Vehicles have battery constraints and charging delays
- Operators must balance service quality vs operational cost



# Project Objectives

- Build a realistic simulation of an urban micromobility system
- Model demand, travel, battery, charging and relocation
- Incorporate stochastic exogenous factors (weather, events)
- Evaluate baseline, heuristic, and reinforcement policies learning under multiple scenarios
- Analyze trade-offs between availability, queues and cost

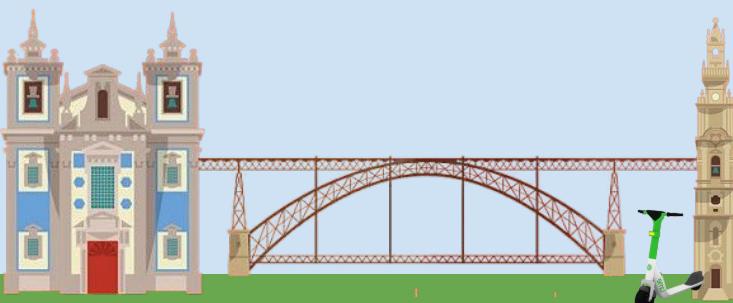


# Methodology



## Space Definition:

- **10 well-spaced station locations** from Porto's bike network (using OpenStreetMap API)
- **Pairwise distances** and **travel times** computed using shortest bike paths
- Defined station parameters: **capacity**, **demand**, **energy consumption**



## Modelling:

- **Discrete-time simulation:** 5-min ticks (steps)
- **Configuration and State**
- **Agent-Environment interface:** RL-style API
  - reset(), step()
  - observations
  - actions ( $act \in \mathbb{R}^4$ )
- **Scenarios and Events:**
  - **Weather:** Discrete-time Markov chain
  - **Demand Peaks:** Morning and evening rush hour
  - **Events:** Random temporary boosts in demand (concerts, sports games, etc.)

# Methodology

**Planners:** Stored in a registry for easy swapping

- **Relocation**

- **Greedy:** From overfilled stations to near-empty ones
- **Budgeted:** Greedy but with a maximum limit per tick

- **Charging**

- **Greedy:** Prioritize charging at low-battery stations
- **Slack:** Charge at stations with low expected demand

*Extra:* **No-op planner** for comparison

## Metrics and KPIs:

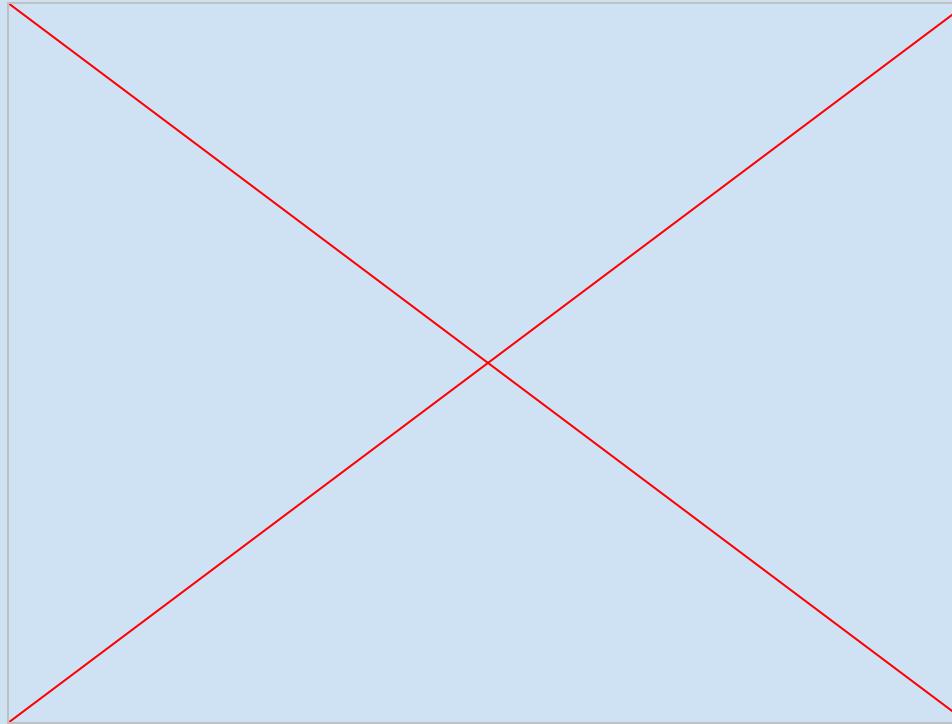


```
1 "ticks": int(T),
2 "dt_min": int(dt_min),
3 # reward + objective
4 "total_reward": float(total_reward),
5 "J_run": float(J_run),
6 "reward_plus_sum": float(reward_plus_sum),
7 "J_avail_run": float(np.mean(J_avail_t)) if T else 0.0,
8 "J_reloc_run": float(np.mean(J_reloc_t)) if T else 0.0,
9 "J_charge_run": float(np.mean(J_charge_t)) if T else 0.0,
10 "J_queue_run": float(np.mean(J_queue_t)) if T else 0.0,
11 # service
12 "demand_total": int(0),
13 "served_total": int(Sunits),
14 "served_new_total": int(snew),
15 "unmet_total": int(U),
16 "availability_tick_avg": safe_mean(availability),
17 "availability_demand_weighted": float(availability_demand_weighted),
18 "unmet_rate": float(unmet_rate),
19 "avg_wait_min_proxy": float(avg_wait_min_proxy),
20 # queue levels + tails
21 "queue_total_avg": safe_mean(queue_total),
22 "queue_total_p95": pct(queue_total, 95),
23 "queue_total_p99": pct(queue_total, 99),
24 "queue_total_max": safe_max(queue_total),
25 "queue_rate_avg": safe_mean(queue_rate),
26 "queue_rate_p95": pct(queue_rate, 95),
27 "queue_rate_p99": pct(queue_rate, 99),
28 "queue_rate_max": safe_max(queue_rate),
29 # SLA time-above-threshold
30 "frac_ticks_queue_gt_10": float(frac_ticks_queue_gt_10),
31 "frac_ticks_queue_gt_20": float(frac_ticks_queue_gt_20),
32 # queue stability
33 "dq_mean": float(dq_mean),
34 "dq_p95": float(dq_p95),
35 "dq_p99": float(dq_p99),
36 "dq_max": float(dq_max),
```



```
1 # operations totals
2 "relocation_km_total": float(reloc_km_total),
3 "reloc_units_total": int(reloc_units.sum()),
4 "reloc_edges_total": int(reloc_edges.sum()),
5 "charging_energy_kwh_total": float(charge_energy_total),
6 "charging_cost_eur_total": float(charge_cost_total),
7 "plugged_total": int(plugged.sum()),
8 "plugged_reserve_total": int(plugged_reserve.sum()),
9 "charge_utilization_avg": safe_mean(charge_util),
10 # operations burstiness (per-tick tails)
11 "reloc_km_p95": pct(reloc_km, 95),
12 "reloc_km_p99": pct(reloc_km, 99),
13 "reloc_units_p95": pct(reloc_units, 95),
14 "reloc_units_p99": pct(reloc_units, 99),
15 # efficiency ratios
16 "km_per_served": float(km_per_served),
17 "eur_per_served": float(eur_per_served),
18 "kwh_per_served": float(kwh_per_served),
19 # system state distribution proxies
20 "empty_ratio_avg": safe_mean(empty_ratio),
21 "empty_ratio_p95": pct(empty_ratio, 95),
22 "empty_ratio_max": safe_max(empty_ratio),
23 "full_ratio_avg": safe_mean(full_ratio),
24 "full_ratio_p95": pct(full_ratio, 95),
25 "full_ratio_max": safe_max(full_ratio),
26 "stock_std_avg": safe_mean(stock_std),
27 "stock_std_p95": pct(stock_std, 95),
28 "fill_p10_avg": safe_mean(fill_p10),
29 "fill_p90_avg": safe_mean(fill_p90),
30 "fill_spread_avg": safe_mean(fill_p90 - fill_p10),
31 # energy / SoC health
32 "soc_mean_avg": safe_mean(soc_mean),
33 "soc_mean_vehicles_avg": safe_mean(soc_mean_vehicles),
34 "soc_station_min_avg": safe_mean(soc_station_min),
35 "soc_station_min_p95": pct(soc_station_min, 5),
36 "soc_station_p10_avg": safe_mean(soc_station_p10),
37 "frac_ticks_soc_p10_lt_0_2": float(frac_ticks_soc_p10_lt_0_2),
38 # rentability primitives (great for discussion)
39 "rentable_frac_avg": safe_mean(rentable_frac),
40 "soc_bind_frac_avg": safe_mean(soc_bind_frac),
41 # overflow
42 "overflow_rerouted_total": int(overflow_rerouted.sum()),
43 "overflow_dropped_total": int(overflow_dropped.sum()),
44 "overflow_extra_min_total": float(overflow_extra_min.sum()),
45 # score weights and scales
46 "score_cfg": asdict(env.score_cfg),
```

# Simulation Demonstration



# Experimental Work: Algorithms



**Heuristic Agent:** Manually crafted rules

- Throttles actions based on previous performance (**reloc. and charging EWMA**)
- Adapts to time-of-day, waiting queues, station occupation

**Episodic Multi-Armed Bandit:** UCB1

- Search of different combinations (arms) of **relocation and charging budgets**
- **One combination** per simulation episode
- After many episodes, the arm with the **best average reward** is learnt (for that specific constant action)

**Contextual Bandit:** LinUCB

- Re-selects budget combinations **every "block"** (e.g. hour), instead of every episode
- **Uses the state of the system** and a **linear model** that predicts the expected reward to pick an arm

**Soft-Actor Critic (SAC):** RL

- Learns a policy (**Deep Neural Network**) that maps **continuous** observations to actions
- Finds highest cumulative reward
- Improves modelling of **uncertainty in demand** (events, weather) through **entropy maximization**

# Experimental Work – Setup



## Chosen KPIs:

- **Availability:** measures service quality
- **Queue Size:** captures user waiting and congestion
- **Relocation Distance:** reflects operational effort
- **Charging cost:** accounts for energy expenditure

## Normalized per-tick objective:

$$J_t = w_A \cdot \frac{U_t}{A_0} + w_R \cdot \frac{R_t}{R_0} + w_C \cdot \frac{C_t}{C_0} + w_Q \cdot \frac{Q_t}{Q_0}$$

## Calibration procedure:

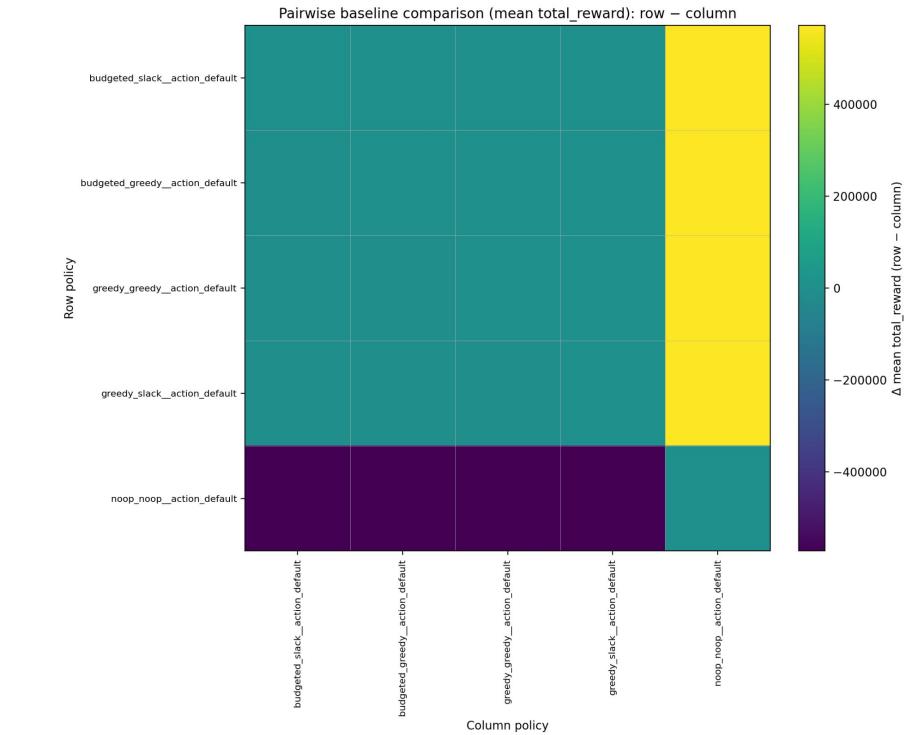
- Each raw signal has different units and scales
- A calibration script was used to
  - Run reference policies under nominal conditions
  - Estimate typical magnitudes of each signal
  - Set normalization constants so that all terms contribute comparably

## Weight design rationale

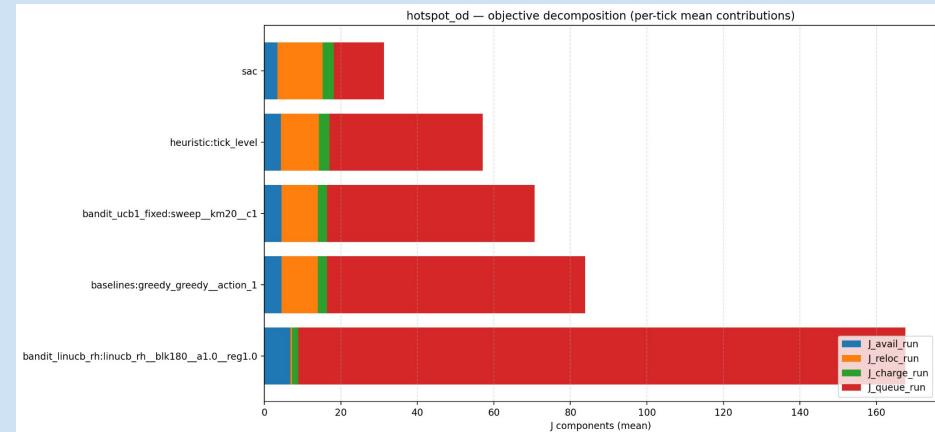
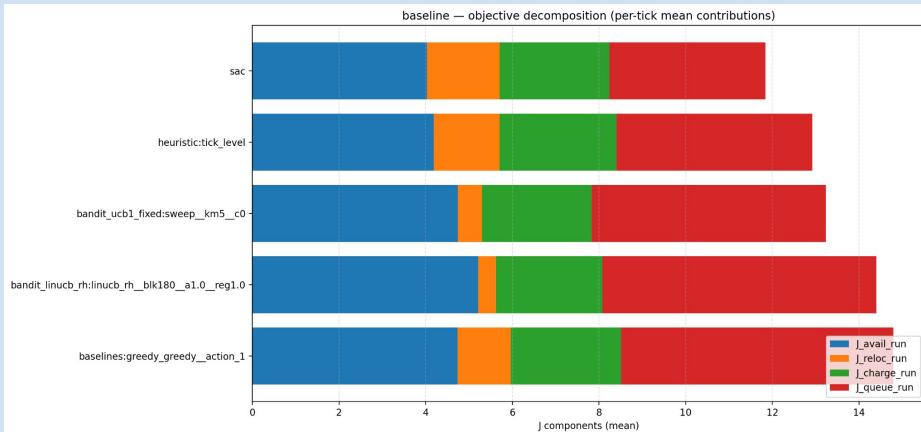
- $w_A = 5$
- $w_Q = 4$
- $w_R = 1$
- $w_C = 1$

# Baselines

- **Baseline set:** no-op, greedy, budgeted, slack
- **Evaluation:** same seeds, same action
- **Metric:** mean total reward (pairwise)
- **Observation:** all active baselines ≫ no-op
- **Observation:** Greedy, budgeted, and slack are near-equivalent in reward



# Analysis of Results



# Critical Analysis and Directions for Future Work



## Critical Analysis

- Performance differences are dominated by queue dynamics, not operational costs
- Queue penalties dominate J\_run in all scenarios
- Heuristics fail under demand heterogeneity
- Bandits reduce cost but sacrifice service quality



## Future Work

- Queue-aware learning objectives
- Multi-timescale decision making:
  - Fast loop for queue mitigation
  - Slow loop for fleet balancing
- Different frequencies so that agents can react to sudden spikes in demand, while not affecting the management in the long run (can be affected by noisy rewards from per tick action).
- Evaluate agents under different maps:
  - Maps with location nodes farther away from each other
  - Maps with more points overall

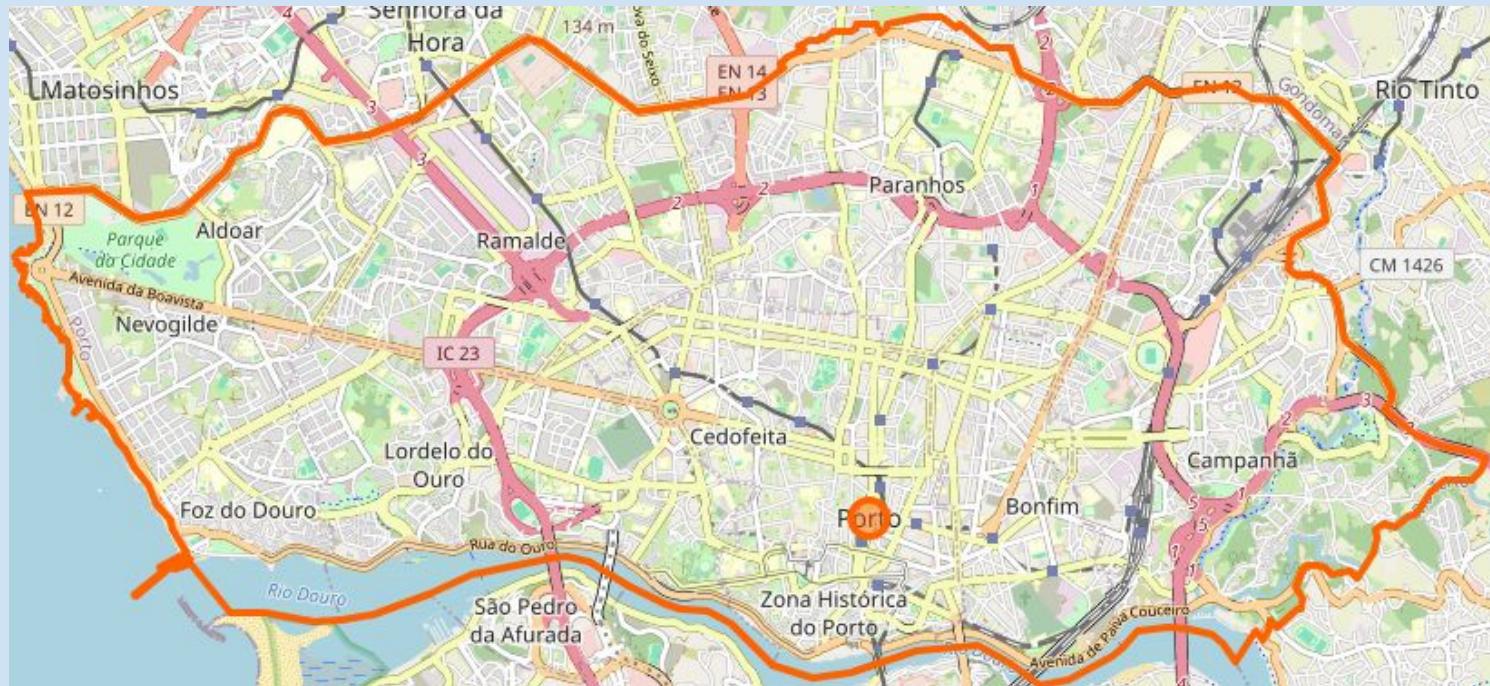


# Thank you!

Any questions?

# Backup slides

# Physical Space Generation



# Modelling: Demand



```
1 """Circular Gaussian over 24h."""
2 h = np.mod(hour, 24.0)
3 d = np.minimum(np.abs(h - mu), 24.0 - np.abs(h - mu))
4 return np.exp(-(d**2) / (2 * sigma**2))
```



```
1 fac = np.ones((horizon_steps, N), dtype=float) # demand factor
2 for _ in range(2):
3     i = int(rng.integers(0, N)) # station index
4     length = int(rng.integers(36, 60)) # ticks if dt=5min (~3-5h)
5     start = int(rng.integers(0, max(1, horizon_steps - length)))
6     fac[start : start + length, i] = 1.8 # +80% demand boost
7 return fac
```



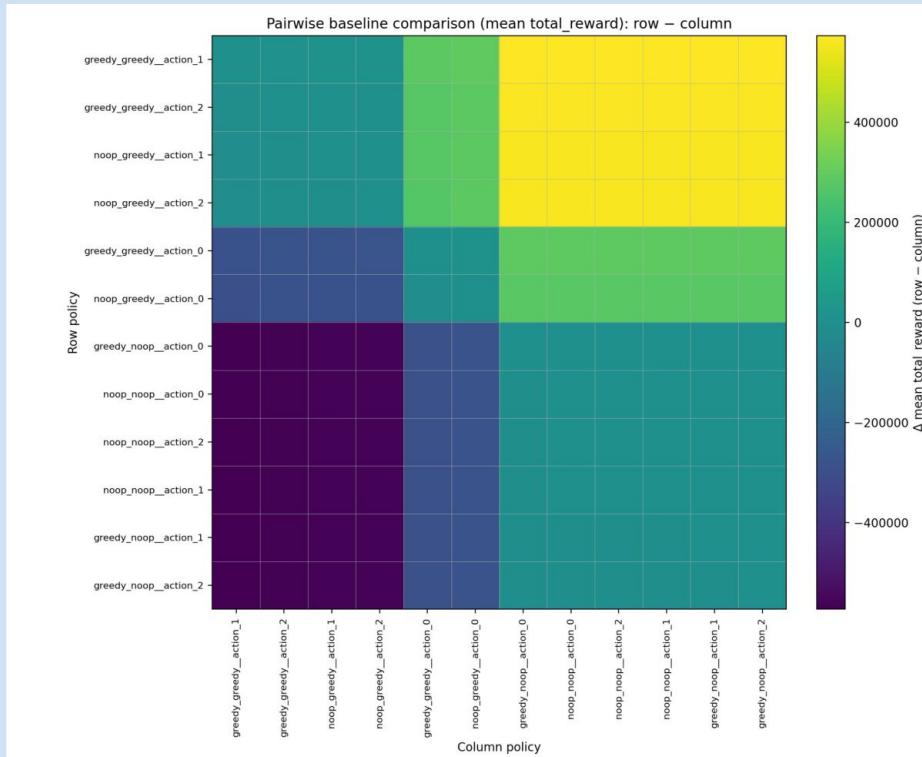
```
1 P = np.array(
2     [
3         [0.86, 0.11, 0.03, 0.00], # clear -> ...
4         [0.10, 0.80, 0.09, 0.01], # cloudy -> ...
5         [0.05, 0.25, 0.65, 0.05], # rain -> ...
6         [0.02, 0.18, 0.45, 0.35], # storm -> ...
7     ],
8     dtype=float,
9 )
10 factors = {"clear": 1.00, "cloudy": 0.90, "rain": 0.60, "storm": 0.45}
```

# Modelling: Full Demand



```
1 """Compute effective demand rate vector at given hour."""
2 lam = diurnal(base_lambda_vec, hour)
3 lam = lam * weather_fac
4 if event_fac_vec is not None:
5     lam = lam * event_fac_vec
6 return lam
```

# Planner Algorithms Comparison



# More Scenarios

